

Decision Defense & Model Ethics

Practical machine learning is about more than achieving high accuracy; it is about justifying why a model is ready for the real world. This session introduces "Decision Defense," focusing on the documentation and ethical standards required to validate results. We will discuss how to identify hidden biases, ensure technical transparency, and build the professional accountability needed to support AI-driven decisions.

The Decision Defense Framework

Decision defense illuminates us and stakeholders about illusion of success (ie. high accuracy).

Stakeholders don't buy "scores"; they buy reliability. Their core question is more like: "Can I trust this model with \$1M or a customer's safety?"

Move from showing what the model found to defending how it reached that conclusion.



The Decision Defense Framework

You will need to defend all actions based on the audit trail established by your ledger and Github history. The idea is to link every experiment to a specific GitHub hash.

What else is important?

- The "Why" Column: Documenting the rationale for feature selection and hyperparameter choices.



The Decision Defense Framework

What else is important? (cnt'd)

- Failure Analysis: A professional defense includes a detailed log of what didn't work and why it was rejected.
- Reproducibility: Ensuring any colleague can recreate your best result using only your documentation.



The Decision Defense Framework

At some point, stress testing & residual integrity become topics all of a sudden. You need to be ready for that moment.

Find where your model breaks:

- Residual Patterns: Proving that your errors are "White Noise." If there is a pattern in your mistakes, your model is incomplete.



The Decision Defense Framework

Find where your model breaks (cnt'd):

- Edge Case Performance: How does the model behave with missing data or extreme outliers?
- Naive Baseline Comparison: Can you outperform a simple "last value" or "average" forecast? If not, the model is overhead, not an asset.



The Decision Defense Framework

Your validation strategy gets to be discussed as well.

Having done segmented validation is a great starting point.

- Global vs. Local Accuracy: A model that is 90% accurate globally might be 0% accurate for a critical sub-group.



The Decision Defense Framework

Having done segmented validation is a great starting point (cnt'd).

- Risk Concentration: Identifying if errors are clustered in high-value or high-risk segments (e.g., fraud detection or premium customers).

A very good defense would present directly (without asking) performance metrics broken down by category, time period, or demographic to prove consistency.



The Decision Defense Framework

Preparing for the Final Review:

- Honest Validation: Did you leak information from the future into the past during training?
- Translation Tax: What is the operational cost of implementing this model versus the current manual process?
- Accountability: Be prepared to explain your model's biggest failures as clearly as its successes.

Is the model "Decision-Ready" or just an "Academic Exercise"?




Model Ethics & The Integrity of Data

In supervised learning, algorithms do not create bias; they inherit it from your data. If your training data contains historical prejudice or systemic errors, your model will codify and automate them.

Question where the labels came from.

- Was the human grader biased?
- Are you creating an echo chamber where models create feedback loops that reinforce existing stereotypes.

MODEL ETHICS & DATA INTEGRITY




Model Ethics & The Integrity of Data

Check out for “proxies”: Even if you remove "Gender" or "Race," variables like "Zip Code" or "University" often act as proxies for the same information.

Find out if there are any “representation gaps”: Does the model fail because the data lacks enough examples of a specific subgroup?


MODEL ETHICS & DATA INTEGRITY



Model Ethics & The Integrity of Data

Impact analysis of representation issues:
Analyzing who is left out when a model
optimizes for the "average" user.

MODEL ETHICS & DATA INTEGRITY




Model Ethics & The Integrity of Data

Are ML models always black boxes?

- Interpretability vs. Accuracy: The professional trade-off. Is a 1% gain in accuracy worth losing the ability to explain a decision?
- Local Explanations: Being able to answer: "Why was this specific application rejected?"
- Stakeholder Rights: The ethical right of an individual to understand an automated decision that affects their life or finances.

MODEL ETHICS & DATA INTEGRITY




Model Ethics & The Integrity of Data

Apply a “clean hands policy” to your data processing process.

- Data Provenance: Knowing the origin, ownership, and "consent status" of every row in your dataset.
- Sovereign Data: Ensuring data residency and security; treating data not as a commodity, but as a strategic asset with ethical boundaries.

MODEL ETHICS & DATA INTEGRITY




Model Ethics & The Integrity of Data

Apply a “clean hands policy” to your data processing process (cnt’d).

- Validation Ethics: Proving that you didn't "torture the data until it confessed" by cherry-picking validation sets.

MODEL ETHICS & DATA INTEGRITY




Model Ethics & The Integrity of Data

Have an audit checklist. It may be simple but its presence creates ground for growth.

- Impact Assessment: What is the worst-case scenario if this model is wrong? This question is a must.

- Bias Monitoring: Do we have a plan to monitor for "Fairness Drift" after the model goes live?

MODEL ETHICS & DATA INTEGRITY




Model Ethics & The Integrity of Data

Have an audit checklist. It may be simple but its presence creates ground for growth (cnt'd).

- Human-in-the-Loop: Identifying the specific points where a human expert must be able to overrule the algorithm.
- Warranty: "We have verified that this model does not systematically disadvantage any subgroup for the sake of global optimization."


MODEL ETHICS & DATA INTEGRITY



Sustainability & Long-Term Governance

Remember why models die post-deployment.

- The Static Snapshot
- Data Drift
- Concept Drift
- Decay Curve




Sustainability & Long-Term Governance

Monitor your inference results and all feedbacks, if any.

Define "Model Half Life"—at what point does the model's error cost more than the price of a full retraining cycle?

Treat every model (even variants) as new artifacts (v2.0, v2.1) with their own "Decision Defense" logs.

Test any new model against the old one in.




Sustainability & Long-Term Governance

Ensure the "why" behind the model lives in the organization's documentation (and GitHub), not just in one person's head.


Manage where your model "lives" to ensure long-term data residency and operational autonomy. Not all models need regulations, but when they do they do.

Knowing when a model is no longer salvageable and must be retired or replaced by a fundamentally different architecture.



Sustainability & Long-Term Governance

Real ML isn't about building a statue; it's about managing a living organism.



Thank You



Please follow up on Github and your emails (school emails).

